



Data Analysis

ECE 356
University of Waterloo
Dr. Paul A.S. Ward

Acknowledgment: slides based in part on materials provided by
Silberschatz, Korth, and Sudarshan, copyright 2010
(source: www.db-book.com)



Learning Outcomes

- Decision Support Systems
 - Data Warehousing
 - Data Mining
 - Classification
 - Association Rules
 - Clustering
-
- Textbook sections (6th ed.): Chapter 20



Decision-Support Systems

- **Decision-support systems** are used to make business decisions, often based on data collected by on-line transaction-processing systems.
- Examples of business decisions:
 - Question answering:
 - What items to stock?
 - What insurance premium to change?
 - To whom to send advertisements?
 - Knowledge acquisition
 - Given this data, what insights (that I don't currently know) is it telling me?
- Examples of data used for making decisions
 - Retail sales transaction details
 - Customer profiles (income, age, gender, *etc.*)



Decision-Support Systems: Overview

- **Data analysis** tasks are simplified by specialized tools and SQL extensions
 - Example tasks
 - ▶ For each product category and each region, what were the total sales in the last quarter and how do they compare with the same quarter last year
 - ▶ As above, for each product category and each customer category
- **Statistical analysis** packages (e.g., : S++; R, Weka) can be interfaced with databases
 - Statistical analysis is a large field, but not covered here
- **Data mining** seeks to discover knowledge automatically in the form of statistical rules and patterns from large databases.
- A **data warehouse** archives information gathered from multiple sources, and stores it under a unified schema, at a single site.
 - Important for large businesses that generate data from multiple divisions, possibly at multiple sites
 - Data may also be purchased externally
 - ▶ e.g., Google maps




Data Warehousing

- Data sources often store only current data, not historical data
- Corporate decision making requires a unified view of all organizational data, including historical data
- A **data warehouse** is a repository (archive) of information gathered from multiple sources, stored under a unified schema, at a single site
 - Greatly simplifies querying, permits study of historical trends
 - Shifts decision-support query load away from transaction processing systems



Data Warehousing





Design Issues

- *When and how to gather data*
 - **Source-driven architecture**: data sources transmit new information to warehouse, either continuously or periodically (e.g., at night)
 - **Destination-driven architecture**: warehouse periodically requests new information from data sources
 - Keeping warehouse exactly synchronized with data sources (e.g., using two-phase commit) is too expensive
 - ▶ Usually OK to have slightly out-of-date data at warehouse
 - ▶ Data/updates are periodically downloaded from online transaction processing (OLTP) systems.
- *What schema to use*
 - Schema integration



More Warehouse Design Issues

■ *Data cleansing*

- *E.g.*, correct mistakes in addresses (misspellings, zip code errors)
- If sampling, ensure sample consistency
- **Merge** address lists from different sources and **purge** duplicates

■ *How to propagate updates*

- Warehouse schema may be a (materialized) view of schema from data sources

■ *What data to summarize*

- Raw data may be too large to store on-line
 - Not anymore
 - Walmart: ~5 billion purchases per year in US
 - Shouldn't be a problem these days
- Aggregate values (totals/subtotals) often suffice
- Queries on raw data can often be transformed by query optimizer to use aggregate values

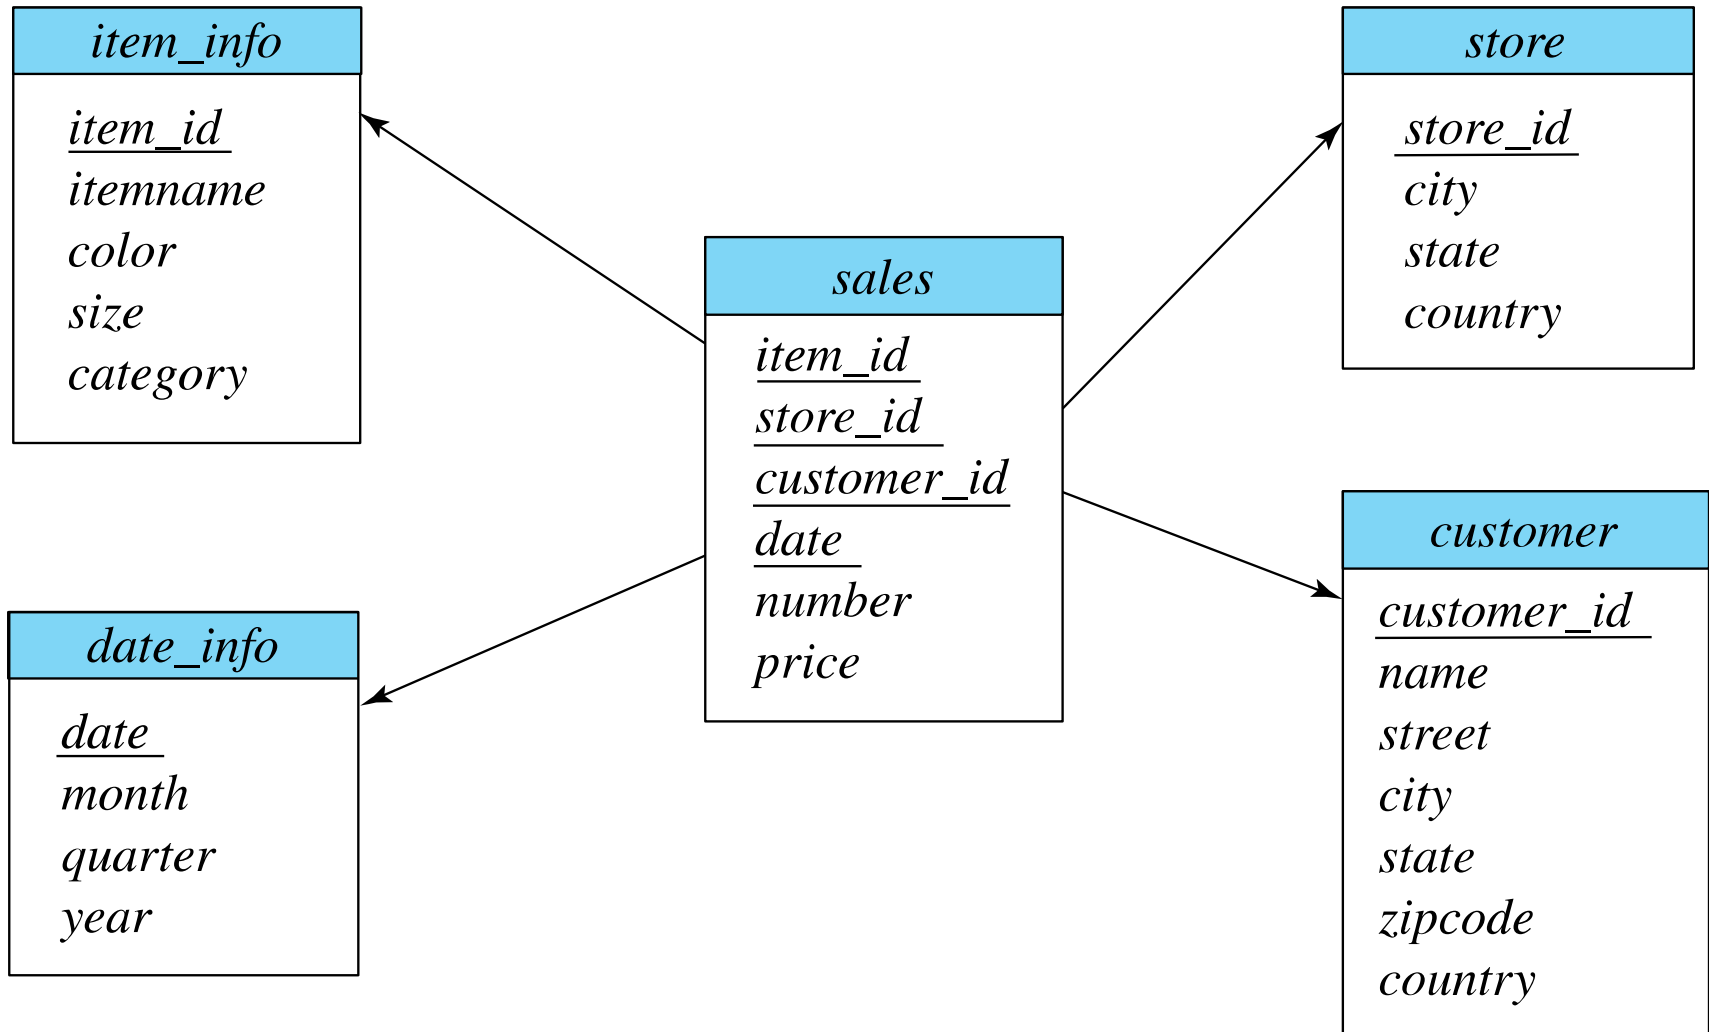


Warehouse Schemas

- Dimension values are usually encoded using small integers and mapped to full values via dimension tables
- Resultant schema is called a **star schema**
 - More complicated schema structures
 - ▶ **Snowflake schema**: multiple levels of dimension tables
 - ▶ **Constellation**: multiple fact tables



Data Warehouse Schema





Data Mining

- Data mining is the process of semi-automatically analyzing large databases to find useful patterns
- **Prediction** based on past history
 - Predict if a credit-card applicant poses a good credit risk, based on some attributes (income, job type, age, ...) and past history
 - Predict if a pattern of credit-card usage is likely to be fraudulent
- Some examples of prediction mechanisms:
 - **Classification**
 - ▶ Given a new item whose class is unknown, predict to which class it belongs
 - **Regression** formulæ
 - ▶ Given a set of mappings for an unknown function, predict the function result for a new parameter value



Data Mining (Cont.)

■ Descriptive Patterns

● Associations

- ▶ Find books that are often bought by “similar” customers. If a new such customer buys one such book, suggest the others too.

● Associations may be used as a first step in detecting **causation**

- ▶ *E.g.*, association between exposure to chemical X and cancer

● Clusters

- ▶ *E.g.*, typhoid cases were clustered in an area surrounding a contaminated well
- ▶ Detection of clusters remains important in detecting epidemics

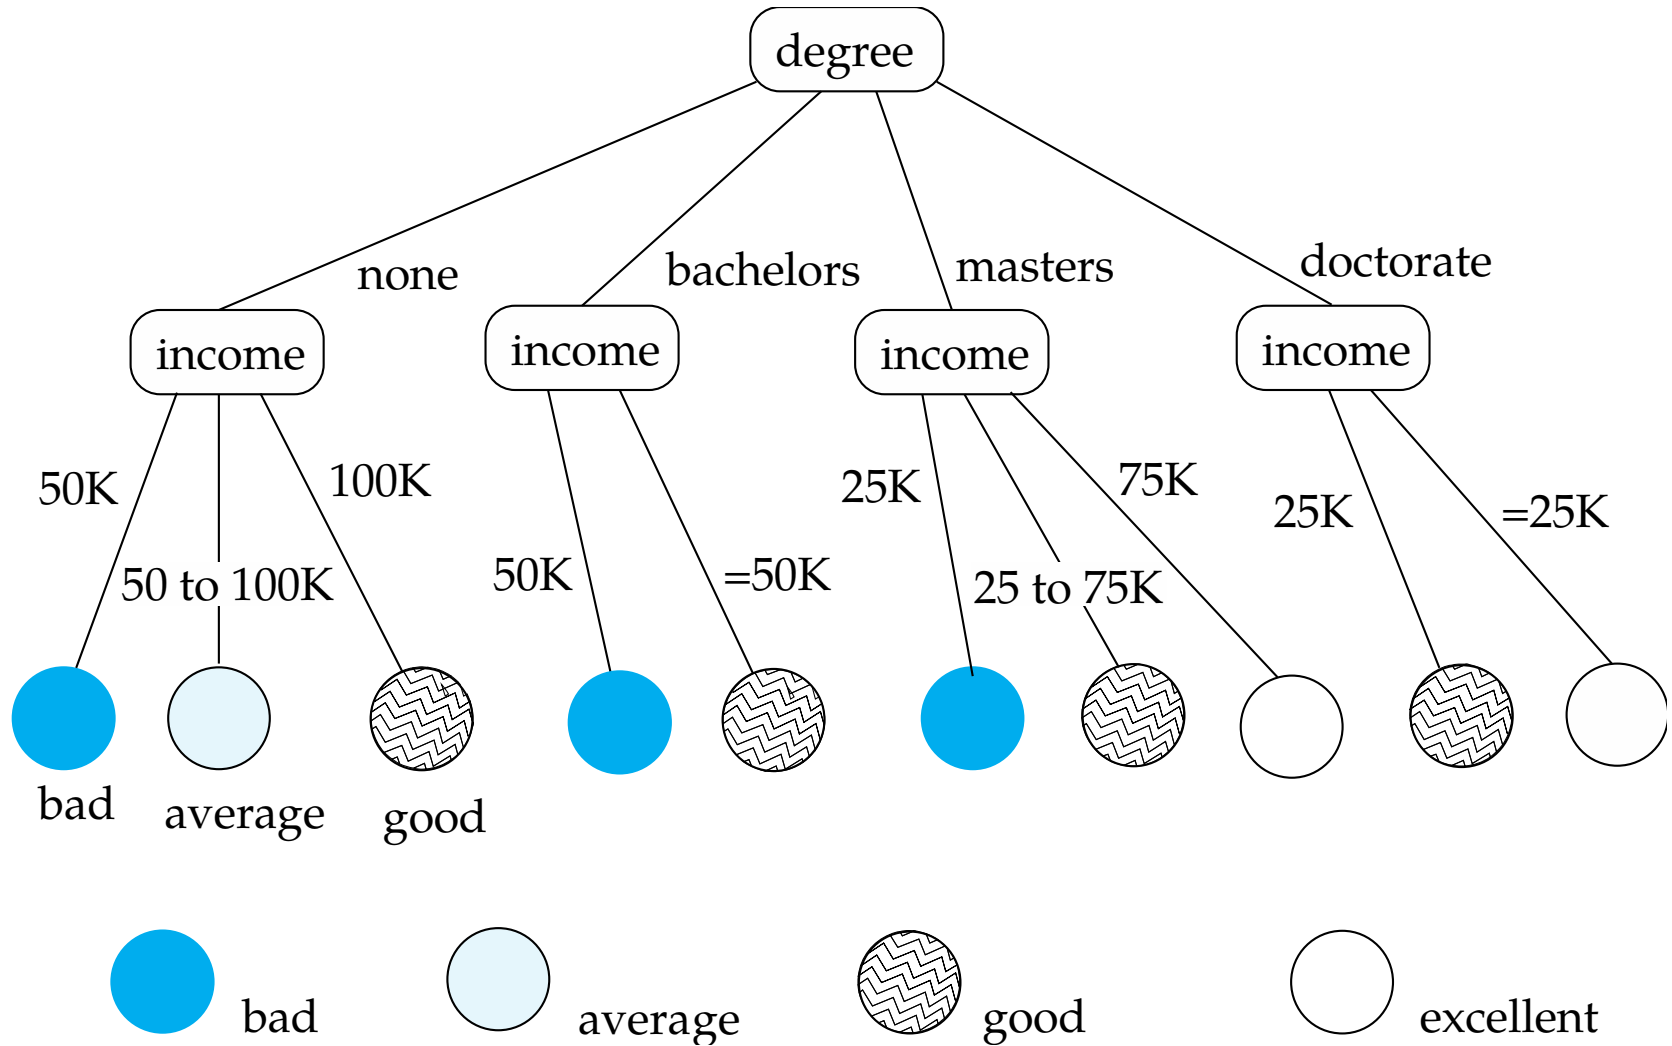


Classification Rules

- Classification rules help assign new objects to classes.
 - *E.g.*, given a new automobile insurance applicant, should he or she be classified as low risk, medium risk or high risk?
- Classification rules for above example could use a variety of data, such as educational level, salary, age, *etc.*
 - $\forall \text{person } P, P.\text{degree} = \text{masters} \textbf{ and } P.\text{income} > 75,000$
 $\rightarrow P.\text{credit} = \text{excellent}$
 - $\forall \text{person } P, P.\text{degree} = \text{bachelors} \textbf{ and }$
 $(P.\text{income} > 25,000 \text{ and } P.\text{income} < 75,000)$
 $\rightarrow P.\text{credit} = \text{good}$
- Rules are not necessarily exact: there may be some misclassifications
- Classification rules can be shown compactly as a decision tree.



Decision Tree





Construction of Decision Trees

- **Training set:** a data sample in which the classification is already known.
- **Greedy** top down generation of decision trees.
 - Each internal node of the tree partitions the data into groups based on a **partitioning attribute**, and a **partitioning condition** for the node
 - **Leaf** node:
 - ▶ all (or sufficient) of the items at the node belong to the same class, **or**
 - ▶ all attributes have been considered, and no further partitioning is possible.



Best Splits

- Pick best attributes and conditions on which to partition
- The impurity of a set S of training instances can be measured quantitatively in several ways.
 - Notation: number of classes = k , number of instances = $|S|$, fraction of instances in class $i = p_i$.
- The **Gini** measure of impurity is defined as:

$$\text{Impurity}(S) = \text{Gini}(S) = 1 - \sum_{i=1}^k p_i^2$$

- When all instances are in a single class, the Gini value is 0
- It reaches its maximum (of $1 - \frac{1}{k}$) if each class has the same number of instances.



Best Splits (2)

- Another measure of impurity is the **entropy** measure, which is defined as:

$$\text{Impurity}(S) = \text{Entropy}(S) = - \sum_{i=1}^k p_i \log_2 p_i$$

- A third measure of purity is **Classification Error** which is:

$$\text{Impurity}(S) = \text{Classification Error}(S) = 1 - \max (p_i)$$



Best Splits (3)



When a set S is split into multiple sets S_i , $i=1, 2, \dots, r$, we can measure the impurity of the resultant set of sets as:

$$\text{Impurity}(S_1, S_2, S_3, \dots, S_r) = \sum_{i=1}^r \frac{|S_i|}{|S|} \text{Impurity}(S_i)$$



The information gain due to particular split of S into S_i , $i = 1, 2, \dots, r$:

InformationGain ($S, \{S_1, S_2, \dots, S_r\}$) = $\text{Impurity}(S) - \text{Impurity}(S_1, S_2, \dots, S_r)$



Best Splits (4)

- Measure of “cost” of a split:

$$\text{InformationContent}(S, S, \{S_1, S_2, \dots, S_r\}) = - \sum_{i=1}^r \frac{|S_i|}{|S|} \log_2 \frac{|S_i|}{|S|}$$

- **Information-gain ratio:**

$$\text{InformationGainRatio} = \frac{\text{InformationGain}(S, S, \{S_1, S_2, \dots, S_r\})}{\text{InformationContent}(S, S, \{S_1, S_2, \dots, S_r\})}$$

- The best split is the one that gives the maximum information-gain ratio



Finding Best Splits

- Categorical attributes (with no meaningful order):
 - Multi-way split, one child for each value
 - Binary split: try all possible breakup of values into two sets, and pick the best
- Continuous-valued attributes (can be sorted in a meaningful order)
 - Binary split:
 - ▶ Sort values, try each as a split point
 - *E.g.*, if values are 1, 10, 15, 25,
 - » split at ≤ 1 , ≤ 10 , ≤ 15
 - » Pick the value that gives best split
 - Multi-way split:
 - ▶ A series of binary splits on the same attribute has roughly equivalent effect



Decision-Tree Construction Algorithm

■ **Procedure** *GrowTree* (S)

 Partition (S);

Procedure Partition (S)

if (*impurity* (S) $< \delta_p$ or $|S| < \delta_s$) **then**

return;

for each attribute A

 evaluate splits on attribute A ;

 Use best split found (across all attributes) to partition

S into S_1, S_2, \dots, S_r

for $i = 1, 2, \dots, r$

 Partition (S_i);



Other Types of Classifiers

- Neural net classifiers are studied in artificial intelligence and are not covered here
- Bayesian classifiers use **Bayes theorem**, which says

$$p(c_j | d) = \frac{p(d | c_j) p(c_j)}{p(d)}$$

where

$p(c_j | d)$ = probability of instance d being in class c_j ,

$p(d | c_j)$ = probability of generating instance d given class c_j ,

$p(c_j)$ = probability of occurrence of class c_j , and

$p(d)$ = probability of instance d occurring



Naïve Bayesian Classifiers

- Bayesian classifiers require
 - computation of $p(d | c_j)$
 - precomputation of $p(c_j)$
 - $p(d)$ can be ignored since it is the same for all classes
- To simplify the task, **naïve Bayesian classifiers** assume attributes have independent distributions, and thereby estimate
$$p(d | c_j) = p(d_1 | c_j) * p(d_2 | c_j) * \dots * (p(d_n | c_j))$$
 - Each of the $p(d_i | c_j)$ can be estimated from a histogram on d_i values for each class c_j
 - ▶ the histogram is computed from the training instances
 - Histograms on multiple attributes are more expensive to compute and store



Regression

- Regression deals with the prediction of a value, rather than a class.
 - Given values for a set of variables, X_1, X_2, \dots, X_n , we wish to predict the value of a variable Y .
- One way is to infer coefficients $a_0, a_1, a_1, \dots, a_n$ such that
$$Y = a_0 + a_1 * X_1 + a_2 * X_2 + \dots + a_n * X_n$$
- Finding such a linear polynomial is called **linear regression**.
 - In general, the process of finding a curve that fits the data is also called **curve fitting**.
- The fit may only be approximate
 - because of noise in the data, or
 - because the relationship is not exactly a polynomial
- Regression aims to find coefficients that give the best possible fit.



Association Rules

- Retail shops are often interested in associations between different items that people buy.
 - Someone who buys bread is quite likely also to buy milk
 - A person who bought the book *Database System Concepts* is quite likely also to buy the book *Operating System Concepts*.
- Associations information can be used in several ways.
 - *E.g.*, when a customer buys a particular book, an online shop may suggest associated books.
- **Association rules:**
 - bread* → *milk* *DB-Concepts, OS-Concepts* → *Networks*
 - Left-hand side: **antecedent**, right-hand side: **consequent**
 - An association rule must have an associated **population**; the population consists of a set of **instances**
 - ▶ *E.g.*, each transaction (sale) at a shop is an instance, and the set of all transactions is the population



Association Rules (Cont.)

- Rules have an associated support, as well as an associated confidence.
- **Support** is a measure of what fraction of the population satisfies both the antecedent and the consequent of the rule.
 - *E.g.*, suppose only 0.001 percent of all purchases include milk and screwdrivers. The support for the rule is *milk* \rightarrow *screwdrivers* is low.
- **Confidence** is a measure of how often the consequent is true when the antecedent is true.
 - *E.g.*, the rule *bread* \rightarrow *milk* has a confidence of 80 percent if 80 percent of the purchases that include bread also include milk.



Finding Association Rules

- We are generally only interested in association rules with reasonably high support (e.g., support of 2% or greater)
- Naïve algorithm
 1. Consider all possible sets of relevant items.
 2. For each set find its support (i.e., count how many transactions purchase all items in the set).
 - ★ **Large itemsets**: sets with sufficiently high support
 3. Use large itemsets to generate association rules.
 1. From itemset A generate the rule $A - \{b\} \rightarrow b$ for each $b \in A$.
 - ✓ Support of rule = support (A).
 - ✓ Confidence of rule = support (A) / support ($A - \{b\}$)



Finding Support

- Determine support of itemsets via a single pass on set of transactions
 - Large itemsets: sets with a high count at the end of the pass
- If memory not enough to hold all counts for all itemsets use multiple passes, considering only some itemsets in each pass.
- Optimization: Once an itemset is eliminated because its count (support) is too small none of its supersets needs to be considered.
- The ***a priori*** technique to find large itemsets:
 - Pass 1: count support of all sets with just 1 item. Eliminate those items with low support
 - Pass i : **candidates**: every set of i items such that all its $i-1$ item subsets are large
 - ▶ Count support of all candidates
 - ▶ Stop if there are no candidates



Other Types of Associations

- Basic association rules have several limitations
- Deviations from the expected probability are often more interesting
 - *E.g.*, if many people purchase bread, and many people purchase cereal, quite a few would be expected to purchase both
 - We are interested in **positive** as well as **negative correlations** between sets of items
 - ▶ Positive correlation: co-occurrence is higher than predicted
 - ▶ Negative correlation: co-occurrence is lower than predicted
- Sequence associations / correlations
 - *E.g.*, whenever bonds go up, stock prices go down in 2 days
- Deviations from temporal patterns
 - *E.g.*, deviation from a steady growth
 - *E.g.*, sales of winter wear go down in summer
 - ▶ Not surprising, part of a known pattern.
 - ▶ Look for deviation from value predicted using past patterns




Clustering

- Clustering: Intuitively, finding clusters of points in the given data such that similar points lie in the same cluster
- Can be formalized using distance metrics in several ways
 - Group points into k sets (for a given k) such that the average distance of points from the centroid of their assigned group is minimized
 - ▶ Centroid: point defined by taking average of coordinates in each dimension.
 - Another metric: minimize average distance between every pair of points in a cluster
- Has been studied extensively in statistics, but on small data sets
 - Data mining systems aim at clustering techniques that can handle very large data sets
 - *E.g.*, the Birch clustering algorithm (slide 32)



Hierarchical Clustering

- Example from biological classification
 - (the word classification here does not mean a prediction mechanism)



- Other examples: Internet directory systems (e.g., Yahoo, more on this later)
- **Agglomerative clustering algorithms**
 - Build small clusters, then cluster small clusters into bigger clusters, and so on
- **Divisive clustering algorithms**
 - Start with all items in a single cluster, repeatedly refine (break) clusters into smaller ones



Clustering Algorithms

- Clustering algorithms have been designed to handle very large datasets
- *E.g.*, the **Birch algorithm**
 - Main idea: use an in-memory R-tree to store points that are being clustered
 - Insert points one at a time into the R-tree, merging a new point with an existing cluster if it is less than some δ distance away
 - If there are more leaf nodes than fit in memory, merge existing clusters that are close to each other
 - At the end of first pass we get a large number of clusters at the leaves of the R-tree
 - ▶ Merge clusters to reduce the number of clusters



Collaborative Filtering

- Goal: predict what movies/books/... a person may be interested in, on the basis of
 - Past preferences of the person
 - Other people with similar past preferences
 - The preferences of such people for a new movie/book/...
- One approach based on repeated clustering
 - Cluster people on the basis of preferences for movies
 - Then cluster movies on the basis of being liked by the same clusters of people
 - Again cluster people based on their preferences for (the newly created clusters of) movies
 - Repeat above till equilibrium
- Above problem is an instance of **collaborative filtering**, where users collaborate in the task of filtering information to find information of interest



Other Types of Mining

- **Text mining:** application of data mining to textual documents
 - cluster Web pages to find related pages
 - cluster pages a user has visited to organize their visit history
 - classify Web pages automatically into a Web directory
- **Data visualization** systems help users examine large volumes of data and detect patterns visually
 - Can visually encode large amounts of information on a single screen
 - Humans are very good at detecting visual patterns